

1 Introduction

The purpose of this paper is to explore Markov Decision Processes. Specifically, analyzing algorithms for solving MDPs in two different problem domains. Those domains will have differing characteristics such as state domain, size, and action types in those domains. I will attempt to compare the results from three different algorithms, Value Iteration, Policy Iteration, and Q-Learning.

2 Problem Descriptions

In this section I will talk about the different problems I will be analyzing and why they are interesting.

2.1 Small Grid World

This problem is fairly classic in reinforcement learning algorithm analysis. The domain of this problem is a 10 x 10 matrix of squares representing a grid which is a subset of the world. (See figure 1 for an example grid world.) There are up to 100 unique states in this world and 4 actions that can be taken in any state. If we suppose that the grid world were a map with the upward direction being North and the remaining three directions, then moving in one of these four directions represent the set of possible actions. Associated with each action there are a few parameters I configured:

- Success rate which controls some stochastic actions into the agents actions in the domain. An action taken by the agent will succeed with probability equal to the success rate $\in [0, 1]$ and do some other random action otherwise.
- Goal reward which controls the reward value of the goal state
- Move reward which controls the value for selecting any action and moving to another square on the grid

Outside of standard walls that bound the outside edge of the grid world I have the option of turning some of the grid squares into "walls" themselves in that they cannot be moved into. This reduces the total number of states, but can actually make traversing the map harder due to the fact that direct routes may not exist. I have done this in my grid world problem to exercise the ability of an algorithm (given discount factor) to consider other potential paths to the goal state.

2.1.1 Why is this problem interesting?

This problem is interesting because it exercises many of the core functionalities of the learning and planning algorithms that I am employing. Due to its low number of states

and actions, as well as its relative simplicity, this problem is compelling in that we should be able to quickly derive insights from and still be able to compare to a more complex one. Since there is also an element of randomness in the actions taken by an agent, this problem is still reasonably complex and requires some form of exploration before directly arriving at an optimal value/policy/Q. The introduction of non-traversable squares should exercise the algorithms ability to consider other possible paths (even those moving in a direction further from the goal.)

2.2 Mountain Car

The mountain car problem is also an interesting single agent reinforcement learning MDP problem. The problem uses a cosine curve to model a valley and two peaks and the agent needs to find its way via actions to the higher of the two (or the top of the mountain.) From the description in the documentation:

”In this domain you can change the parameters for min/max position and velocity, the scale of the cosine curve on which the car travels, the force of gravity, acceleration, and the amount of time that elapses between simulation/decision steps.” [1]

The actions that an agent can take in this domain are:

- Move Backwards
- Move Forwards
- Do Nothing

2.2.1 Why is this problem interesting?

This problem is interesting given that it has a large number of potential states. By large, I mean that it has an infinite number of states given that it is a continuous domain modeled by the cosine curve. In order to overcome the requirement by value iteration and policy iteration in BURLAP to be able to determine all states that are reachable I will be sampling 5000 states from a uniform distribution positions along this curve. Whereas in the grid world there was some stochastic randomness to the transition function between states, the agent in this world will have different external forces to overcome, namely the force of gravity and the lack of friction whereby an agent could perform no actions and still experience a transition (otherwise known as momentum.) Given all these it will be interesting to see the results of the algorithms on this problem.

3 Algorithms

In this section we will talk about the algorithms I will be using to analyze these problems as well as the parameters that I will vary in search of an optimal solution to these problems.

3.1 Parameter Exploration

- γ which is the discount factor where $\gamma \in [0, 1]$ and manages the explore/exploit dilemma in RL where high values favor long-term rewards, and low values optimize for greedier short-term rewards
- δ which controls the threshold that if Policy Iteration, Value Iteration, or Q-Learning improve by less than this value, then the algorithm terminates
- Initial Q Value initialized optimistically in my case of 0.3
- α which controls the learning rate of the respective algorithms
- Number of Trials, which dictates the number of times to re-run a learning session with a particular algorithm, the average reward of which is taken as the final result

Unless explicitly mentioned all algorithms search in an ϵ -greedy fashion, where $\epsilon = 0.1$

3.2 Value Iteration

For both problems I used the BURLAP implementation of value iteration to derive the optimal reward values for each state according to the Bellman update rule. Even though Value Iteration is guaranteed to converge, in order to keep my experiments from taking forever, I leveraged from the params above to terminate the updates after a threshold or maximum iterations occurs.

3.3 Policy Iteration

For both problems I used the BURLAP implementation for Policy iteration passing in from the above section the δ , and max iterations to control the termination provided that we never improve less than δ in a Learning episode.

3.4 Q-Learning

Much like the previous two sections, I leveraged the BURLAP implementation of Q-Learning using the params from above which apply to Q-Learning and optimized them for each problem.

4 Analysis

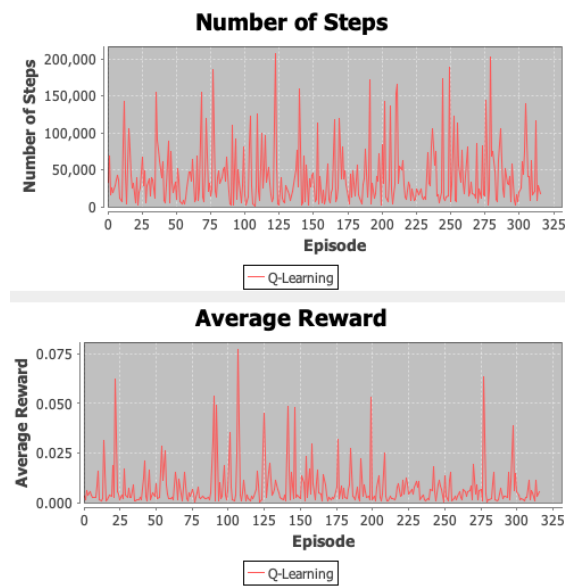
4.1 Small Grid World

Results of algorithms: Due to the number of states, everything listed below is the average of 10 repeated trials as opposed to 100 for Small Grid World.

	Policy Iteration	Value Iteration	Q-Learning
Average Number of Actions	164	37530	47624
Average Reward	100	100	100
Best Number of Actions	164	1683	18646
Best Reward	100	100	100
Average Processing Time	20503.0(s)	118.1(s)	285.9(s)

- How many iterations does it take to converge
- Which one converges faster and why?
- How do I define convergence (delta)
- Do they converge to the same answer
- How did the number of states affect things
- How does it perform especially in comparison to the cases where you knew the model, rewards, and so on?
- What explorations strategies did you choose?
- Did some work better than others?

5 Conclusion



In this paper I have explored the tradeoffs between several popular MDP solving algorithms in BURLAP. I have compared them in two different novel problems which exercised and demonstrated the differences and similarities in terms of performance of these algorithms.

Figure 10: Q-Learning Deep Dive Mountain Car




Figure 1: Small Grid World




Figure 2: Policy Iteration Small Grid World




Figure 3: Value Iteration Small Grid World



Figure 4: Q-Learning Small Grid World

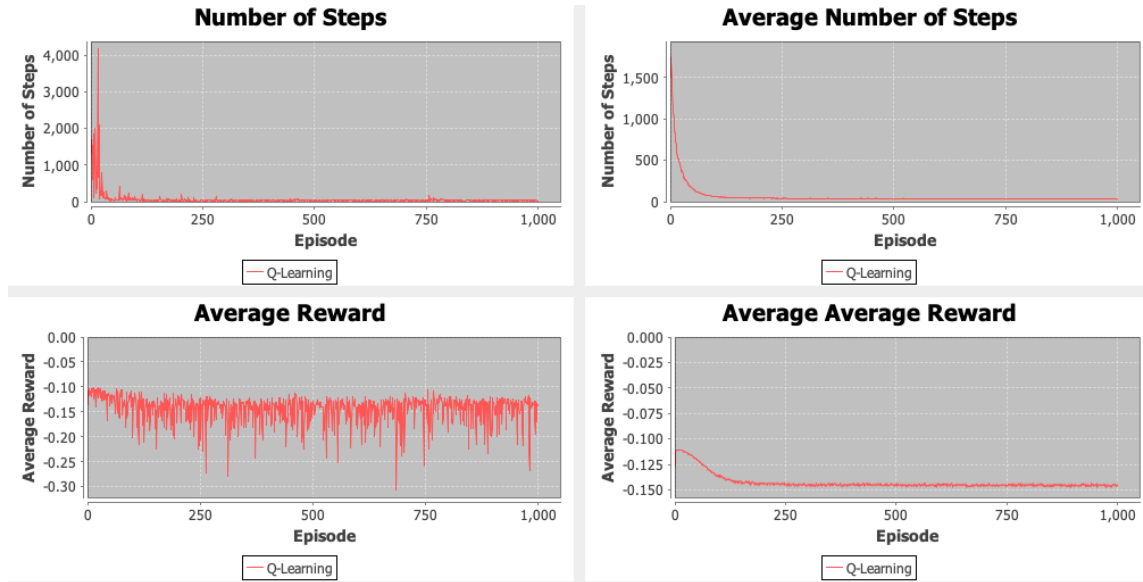


Figure 5: Q Learning Analysis Small Grid World



Figure 6: Value Iteration Mountain Car



Figure 7: Policy Iteration Small Grid World




Figure 8: Mountain Car Domain



Figure 9: Q-Learning Mountain Car

References

- [1] James MacGlashan. Burlap.
<http://burlap.cs.brown.edu/faq.html#cite>, 2020.